

PredictEvent

Event-Driven Congestion Intelligence for Bengaluru Traffic Police


Theme 2 | BTP × Flipkart Hackathon | Forecast impact → recommend manpower, barricading & diversion

The Problem

- Rallies, festivals, sports, processions, VIP movement, construction & sudden incidents choke Bengaluru's corridors.
- Event impact is not quantified in advance — which corridor, how hard, tomorrow?
- Resource deployment is experience-driven — manpower & barricades decided by gut feel.
- No post-event learning loop — the same mistakes repeat.
- Goal: turn 8,173 historical events into a forward-looking deployment planner.

Our Solution


- **A laptop-deployable decision-support engine (no GPU) with 3 part**
 - Hotspot Load Forecaster → forecast load per corridor for tomorrow
 - Closure-Risk Classifier → barricade & RED/AMBER/GREEN triage
 - OR Optimization Engine → SLA-guaranteed officers + coverage-optim
- **Uses ONLY the provided ASTRAM dataset. Interpretable, deployab**



ML Results — honest, time-based split

Component	Result	Baseline
Hotspot Load Forecaster	R^2 0.536 · MAE 2.19	seasonal-naive R^2 -0.32
Closure-Risk Classifier	ROC-AUC 0.813 · recall 0.76	random PR-AUC 0.072
Duration Predictor	not reliably learnable	(data limitation, see findings)

Operating point: recall-floor — catches $\geq 75\%$ of closures (a missed closure costs BTP more).
Top drivers (permutation importance): event_cause $\gg$ corridor > location > vehicle type.
Deep ST-model: MAE 2.12 vs GBM 2.24 (~5% better, R^2 tied); ensemble R^2 0.55. GBM kept in prod.
Calibrated probabilities: ECE 0.33 $\rightarrow$ 0.02, Brier 0.185 $\rightarrow$ 0.058 — '81% risk' really means ~81%.
Cross-validated (5-fold expanding): hotspot R^2 0.56 $\pm$ 0.14, closure AUC 0.73 $\pm$ 0.05 — spread, not cherry-pick.



OR Optimization Engine — the differentiator



- We replaced heuristic rules with the forecast→optimize pattern used in deployed police systems.
- Manpower = M/M/c queueing (Erlang-C): officer count sized so expected response wait $\leq$ SLA (e.g. 5 min) — a number WITH a service-level guarantee.
- Barricade placement = Maximum Coverage Location Problem (MCLP): place P units to cover max risk×load-weighted demand (PuLP/CBC).
- Risk×load weighting: both high-traffic AND high-risk corridors, per OR location-allocation literature.
- Every number is traceable — no black box. Evidence: SMU IJCAI-2019, Vlahogianni 2019, Miao/Easa 2021.

Data Findings That Build Trust

- Timestamps are local IST mislabeled '+00' — proven via bimodal rush-hour signature; we use stored hour as-is.
- `closed_datetime` is administrative ticket-close, NOT on-ground clearance → exact duration not learnable. We refuse to over-claim it.
- We caught & removed target leakage: end-coordinate flag is filled BECAUSE a closure happened (98% vs 0%) — faked ROC-AUC 1.0.
 - Permutation importance exposed it; dropping it gives the honest 0.81.
- Event-driven causes (`public_event`, `procession`, `VIP`, `protest`) have 40–80% closure rates — the true Theme-2 targets.

Prototype — Live Demo

- Event Simulator — type an event (e.g. IPL match @ Chinnaswamy) → instant risk %, tier, personnel, barricade & diversion.
- Hotspot Forecast + Manpower — pick a day → corridor load + recommended personnel.
- Live Event Triage — every event ranked by closure-risk with action tier + map.
- Analytics — historical patterns + feature importance.



Where It Fits in BTP's Stack

- BATCS optimizes the SIGNALS — adaptive timing at 165 junctions (CoSiCoSt). Not people.
- ASTRAM SEES the city — awareness, 15-min alerts, microsimulation. (Our dataset's source.)
- Neither produces a prescriptive police-deployment plan.
- PredictEvent DECIDES the deployment — turns ASTRAM's event data into an officer count with a service-level guarantee + coverage-optimal placement.
- A prescriptive module on top of the stack BTP already runs — it extends ASTRAM, doesn't compete.

Why It Wins · Roadmap

- Solves BTP's actual #1 pain — event manpower planning, not a tech demo.
- OR-optimized, SLA-guaranteed deployment — not heuristic rules.
- Honest, leakage-checked metrics — survives a technical jury.
- Plugs into ASTRAM; deployable today; provided-data only.
- Roadmap:
 - Live ASTRAM API feed → real clearance timestamps
 - Capacitated city-wide MCLP + IP shift-scheduling (ILS/tabu fast replanning)
 - Mobile officer app + automatic retraining loop